





# TPC—DI Project

 Dates	@December 24, 2024
 Course	INFO-H419
 Status	<input type="checkbox"/>
 Task	

## DATA GENERATION

我不知道為什麼但是只有退到用java8才能run

sf 3 =

AuditTotalRecordsSummaryWriter - TotalRecords for Batch1: 4539962

AuditTotalRecordsSummaryWriter - TotalRecords for Batch2: 19900

AuditTotalRecordsSummaryWriter - TotalRecords for Batch3: 19855

AuditTotalRecordsSummaryWriter - TotalRecords all Batches: 4579717 695794.14  
records/second

Overall time 0h:00m:06s:581ms

Generated 267.1 MiB

Speed 40.6 MiB/s

sf 6 =

AuditTotalRecordsSummaryWriter - TotalRecords for Batch1: 9443149

AuditTotalRecordsSummaryWriter - TotalRecords for Batch2: 40317

AuditTotalRecordsSummaryWriter - TotalRecords for Batch3: 40283

AuditTotalRecordsSummaryWriter - TotalRecords all Batches: 9523749  
1074308.97 records/second

Overall time 0h:00m:08s:867ms

Generated 564.2 MiB

Speed 63.6 MiB/s

sf 12 =

AuditTotalRecordsSummaryWriter - TotalRecords for Batch1: 19242651  
AuditTotalRecordsSummaryWriter - TotalRecords for Batch2: 81117  
AuditTotalRecordsSummaryWriter - TotalRecords for Batch3: 81095  
AuditTotalRecordsSummaryWriter - TotalRecords all Batches: 19404863  
1466621.04 records/second

Overall time 0h:00m:13s:232ms

Generated 1.1 GiB

Speed 87.6 MiB/s

sf 24 =

AuditTotalRecordsSummaryWriter - TotalRecords for Batch1: 38842886  
AuditTotalRecordsSummaryWriter - TotalRecords for Batch2: 162844  
AuditTotalRecordsSummaryWriter - TotalRecords for Batch3: 162628  
AuditTotalRecordsSummaryWriter - TotalRecords all Batches: 39168358  
1780218.07 records/second

Overall time 0h:00m:22s:001ms

Generated 2.3 GiB

Speed 107.0 MiB/s

以下都是手動ETL只針對sf3的的code，airflow automate的時候注意看path

## CREATE DATABASE AND SCHEMAS

```
MariaDB [(none)]> CREATE DATABASE tpc_di;
```

```
MariaDB [(none)]> USE tpc_di;
```

```
MariaDB [tpc_di]> CREATE SCHEMA staging;
```

```
MariaDB [tpc_di]> CREATE SCHEMA master;
```

# DESIGN TABLES

```
MariaDB [master]> SHOW TABLES;
```

```
+-----+
| Tables_in_master |
+-----+
| audit              |
| dimaccount         |
| dimbroker          |
| dimcompany         |
| dimcustomer        |
| dimdate            |
| dimessages         |
| dimsecurity        |
| dimtime            |
| dimtrade           |
| factcashbalances   |
| factholdings       |
| factmarkethistory  |
| factwatches        |
| financial          |
| industry           |
| prospect           |
| statustype         |
| taxrate            |
| tradetype          |
+-----+
20 rows in set (0.001 sec)
```

```
[MariaDB [staging]> show tables;
```

```
+-----+
| Tables_in_staging |
+-----+
| audit              |
| batchdate          |
| cashtransaction    |
| customermgmt       |
| dailymarket        |
| date               |
| finwire            |
| finwire_cmp        |
| finwire_fin        |
| finwire_sec        |
| holdinghistory     |
| hr                 |
| industry           |
| prospect           |
| statustype         |
| taxrate            |
| time               |
| trade              |
| tradehistory       |
| tradetype          |
| watchhistory       |
+-----+
21 rows in set (0.000 sec)
```

## LOAD DATA INTO STAGING

Error populating staging tables: (1292, "Incorrect datetime value: '2013-03-31 02:18:26' for column `staging.cashtransaction.ct_dts` at row 39188")

The datetime value `2013-03-31 02:18:26` might fall into a range that is invalid due to Daylight Saving Time (DST) changes or input formatting. For example, in many regions, `2013-03-31 02:18:26` would be invalid because clocks jumped from `01:59:59` to `03:00:00`.

MariaDB may reject certain datetime values if they fall into invalid ranges due to DST. Address this by:

```
SET GLOBAL sql_mode = '';
SET SESSION sql_mode = '';
```

run `load_staging.py` + `populate_staging_customermgmt.py`

## TRANSFORM AND LOAD TO MASTER

1. Historical load phase: Transforms and loads initial data from Batch1 (timed).
2. Incremental update 1: Loads Batch2 data incrementally (timed).
3. Incremental update 2: Loads Batch3 data incrementally (timed).
4. Automated audit: Validates results after final update.

我在github看了一圈，沒有任何組曾經做過234，這後面的工作量超級無敵不合理的恐怖，我就默認我們不需要做了。<https://github.com/QasimKhan5x/TPC-DI/blob/main/scripts/incremental.py> 這有個兩千多行的code可以改著用,可以直接在report說我們refer to這個，如果我們打算做的話。。。 (如果我沒有update說明我放棄了 (99%我會放棄) )

**static tables = These tables are loaded during the Historical Load phase and are not modified afterward:**

run 1.load\_master\_static.sql

- **StatusType:** From `StatusType.txt`
- **TaxRate:** From `TaxRate.txt`
- **TradeType:** From `TradeType.txt`
- **DimDate:** From `Date.txt`
- **DimTime:** From `Time.txt`
- **Industry:** From `Industry.txt`

**non-static tables = These tables are updated dynamically through history tracking or incremental updates:**

- **DimCustomer:** Updated based on incoming data and history tracking mechanisms.
- **DimAccount:** Implements history tracking updates.
- **DimBroker:** History tracked but rarely updated in the benchmark.
- **DimSecurity:** Obtains updates from FINWIRE files through history tracking.
- **DimCompany:** Updated as needed from FINWIRE files.

- **Fact Tables:** Such as `FactHoldings` , `FactMarketHistory` , `FactCashBalances` , and `FactWatches` are regularly updated based on transactions and other inputs

#### other Intermediate tables:

- **Prospect**
  - Temporary staging table for potential customer data.
  - Transforms data into `DimCustomer` through periodic updates.
- **DimMessages**
  - Logging table for system events.
  - Stores ETL process messages and alerts.
- **financial**
  - serves as a bridge between staging data ( `staging.finwire_fin` ) and fact tables ( `factmarkethistory` )

run all 2-18 sql files in Historical Transformation file.

#### Ordered Execution Steps:

##### 1. Load static tables:

##### 2. Load `master.dimcompany` :

- Required for `dimsecurity` , `financial` , `factmarkethistory` , and `dimessages` `dimcompany` .
- **Records: 4500**
- **manually: 0.258 sec**

##### 3. Load `master.dimessages dimcompany` :

- Uses `dimcompany` .
- **Records: 450**
- **manually: 0.010 sec**

##### 4. Transform and load `master.dimbroker` :

- Depends on `dimdate` .

- **Records: 4293**
  - **manually: 0.044 sec**
5. **Transform and load** `master.prospect` :
- Depends on `dimdate` .
  - **Records: 14981**
  - **manually: 1 min 22.277 sec**
6. **Transform and load** `master.dimcustomer` :
- Depends on `taxrate` and `prospect` .
  - **Records: 6440**
  - **manually: 1.223 sec**
7. **Load** `master.dimessages dimcustomer` :
- Depends on `dimcustomer` .
  - **Records: 1720**
  - **manually: 0.016 sec**
8. **Update** `master.prospect` :
- Depends on `dimcustomer` .
  - **Rows matched: 3377 Changed: 3377**
  - **manually: 16.741 sec**
9. **Transform and load** `master.dimaccount` :
- Depends on `dimbroker` and `dimcustomer` .
  - **Records: 11261**
  - **manually: 2 min 9.531 sec**
10. **Transform and load** `master.dimsecurity` :
- Depends on `dimcompany` .
  - **Records: 8838**
  - **manually: 6.715 sec**

## 11. Transform and load `master.dimtrade`:

- Depends on `dimaccount`, `dimsecurity`, `statustype`, and `tradetype`.
- **Records: 3721368**
- **manually: 3 min 40.046 sec**
- query optimisation from 1st(takes more than 30mins) to 2nd(index, reduced row processed, avoid full table scan on trade)

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	tt	ALL	NULL	NULL	NULL	NULL	5	
1	PRIMARY	st	ALL	NULL	NULL	NULL	NULL	6	Using join buffer (flat, BNL join)
1	PRIMARY	s	ALL	NULL	NULL	NULL	NULL	8963	Using join buffer (incremental, BNL join)
1	PRIMARY	t	ALL	NULL	NULL	NULL	NULL	11394	Using join buffer (incremental, BNL join)
1	PRIMARY	t	ALL	NULL	NULL	NULL	NULL	391100	Using where; Using join buffer (incremental, BNL join)
1	PRIMARY	<derived3>	ref	key0	key0	8	staging.t.t_id	10	Using where
1	PRIMARY	<derived2>	ref	key0	key0	8	staging.t.t_id	10	
3	DERIVED	th	ALL	NULL	NULL	NULL	NULL	982497	Using temporary
2	DERIVED	th	ALL	NULL	NULL	NULL	NULL	982497	Using temporary; Using filesort

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	tt	ALL	NULL	NULL	NULL	NULL	5	
1	PRIMARY	st	ALL	NULL	NULL	NULL	NULL	6	Using join buffer (flat, BNL join)
1	PRIMARY	s	ALL	idx_dimsecurity_symbol	NULL	NULL	NULL	8963	Using join buffer (incremental, BNL join)
1	PRIMARY	t	ref	idx_trade_t_ca_id,idx_trade_t_s_symb	idx_trade_t_s_symb	60	master.s.symbol	247	Using where
1	PRIMARY	a	ref	idx_dimaccount_accounttid	idx_dimaccount_accounttid	5	staging.t.t_ca_id	4	
1	PRIMARY	<derived3>	ref	key0	key0	8	staging.t.t_id	10	Using where
1	PRIMARY	<derived2>	ref	key0	key0	8	staging.t.t_id	10	
3	DERIVED	th	ALL	idx_tradehistory_th_t_id,idx_t_id	NULL	NULL	NULL	982497	Using where; Using temporary
2	DERIVED	th	ALL	idx_tradehistory_th_t_id,idx_t_id	NULL	NULL	NULL	982497	Using where; Using temporary; Using filesort

## 12. Load `master.dimessages dimtrade`:

- Depends on `dimtrade`.
- **Records: 312**
- **manually: 2.311 sec**

## 13. Transform and load `master.financial`:

- Depends on `dimcompany`.
- **Records: 496617**
- **manually: 2 min 25.753 sec**

## 14. Transform and load `master.factcashbalances`:

- Depends on `dimaccount` and `dimdate`.
- **Records: 564461**

- manually: 18 min 55.160 sec

#### 15. Transform and load `master.factholdings` :

- Depends on `dimtrade` .
- **Records: 3485814**
- manually: 10.837 sec
- query optimisation from 1st(takes more than 1h) to 2nd(index, avoid full table scan on holdinghistory)

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	h	ALL	NULL	NULL	NULL	NULL	362529	
1	SIMPLE	t	ALL	NULL	NULL	NULL	NULL	3721227	Using where; Using join buffer (flat, BNL join)

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	t	ALL	idx_tradeid	NULL	NULL	NULL	3721311	
1	SIMPLE	h	ref	idx_hh_t_id	idx_hh_t_id	7	master.t.tradeid	1	Using index condition

#### 16. Transform and load `master.factwatches` :

- Depends on `dimcustomer` , `dimsecurity` , and `dimdate` .
- **Records: 1532010**
- manually: 4.846 sec (optimised with indexes on `w_c_id` , `w_s_symb` , `w_action` in `staging.watchhistory` for faster filtering and joins + Indexes on `customerid` , `symbol` , and `datevalue` in related tables)

#### 17. Transform and load `master.factmarkethistory` :

- Depends on `dimdate` , `financial` , `dimcompany` , and `dimsecurity` .
- **Records: 6240582**
- manually: 5 min 42.005 sec (Added indexes to `dailymarket` , `dimsecurity` , `dimdate` , and `financial` for join optimization . tables are filter to only include recent years data because it still takes damn long for mariadb to process)



id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	s	ALL	idx_dimsecurity_symbol	NULL	NULL	NULL	8963	
1	PRIMARY	<derived5>	ref	key0	key0	6	master.s.sk_companyid	55	
1	PRIMARY	<derived4>	ref	key0	key0	63	master.s.symbol	1431	Using where
4	DERIVED	dm	ALL	idx_dailymarket_symbol_date	NULL	NULL	NULL	1282927	Using temporary; Using filesort
4	DERIVED	dd	ref	idx_dimdate_datevalue	idx_dimdate_datevalue	3	staging.dm.dm_date	1	
4	DERIVED	<derived3>	ref	key0	key0	63	staging.dm.dm_v_symb	10	Using where
3	DERIVED	dm	ALL	NULL	NULL	NULL	NULL	1282927	Using temporary
3	DERIVED	dd	ref	idx_dimdate_datevalue	idx_dimdate_datevalue	3	staging.dm.dm_date	1	Using index
5	DERIVED	f	ALL	idx_financial_company_qtr	NULL	NULL	NULL	496713	Using temporary
6	DEPENDENT SUBQUERY	f2	ref	idx_financial_company_qtr	idx_financial_company_qtr	5	master.f.sk_companyid	81	Using where; Using index

## 18. Load `master.dimessages factmarkethistory` :

- Depends on `factmarkethistory` .
- **Records: 6642**
- **manually: 4.231 sec**

"The performance metric reported for TPC-DI is a throughput measure, the number of **Source**

### Data

rows processed per second. Conceptually, it is calculated by dividing the total rows processed by the elapsed time of the run.

The completion timestamp (CT) for each phase is found in `Dlmessages.MessageDateAndTime` where `MessageType='PCR'` and `BatchID` matches the phase. Example query for **Historical Load**:

`select MessageDateAndTime from Dlmessages where BatchID = 1 and MessageType = 'PCR'`

For **Historical Load** throughput calculation:

1. Elapsed time (EH) = CT1 - CT0
2. Total rows (RH) = Batch1 row count from **DIGen**
3. Throughput (TH) = RH / EH"

TPC-DI specification是這麼寫的，但是我看到的組也都只看airflow measure execution time ?

過程referred to: <https://github.com/risg99/tpc-di-benchmark/tree/main> airflow的code改改應該就能用了